

COMP9313: Big Data Management

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Assignment Project Exam Help

Ch <https://eduassistpro.github.io/> **ce l**

Add WeChat edu_assist_pro

What is MapReduce

- ❑ Origin from Google, [OSDI'04]
 - ❑ MapReduce: Simplified Data Processing on Large Clusters
 - ❑ Jeffrey Dean and Sanjay Ghemawat
- ❑ Programming model for parallel data processing
- ❑ Hadoop can run MapReduce programs written in various languages: e.g. Java, Ruby,
- ❑ For large-scale
 - ❑ Exploits large set of commodity
 - ❑ Executes process in distributed
 - ❑ Offers high availability

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Motivation for MapReduce

- A Google server room:

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<https://www.youtube.com/watch?t=3&v=avP5d16wEp0>

Motivation for MapReduce

- Typical big data problem challenges:
 - How do we break up a large problem into smaller tasks that can be executed in **parallel**?
 - How do we assign tasks to workers distributed across a potentially **large number** of machines?
 - How do we ensure that the workers get the **data** they need?
 - How do we **share** the different workers?
 - How do we **share** partial results or that is needed by another?
 - How do we accomplish all of the above in the face of software **errors** and hardware **faults**?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Motivation for MapReduce

- There was need for an abstraction that hides many system-level details from the programmer.
- MapReduce addresses this challenge by providing a simple abstraction for the developer, transparently handling most of the details behind the scenes in a **scalable**, **robust**, and **efficient** manner.

<https://eduassistpro.github.io/>

- MapReduce separates the **what** from the **how**

Add WeChat edu_assist_pro

Jeffrey (Jeff) Dean

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- He is currently a Google Senior Fellow in the Systems and Infrastructure Group
- Designed MapReduce, BigTable, etc.
- One of the most genius engineer, programmer, computer scientist...
- Google “Who is Jeff Dean” and “Jeff Dean facts”

Jeff Dean Facts

- Kenton Varda created "Jeff Dean Facts" as a Google-internal April Fool's joke in 2007.
- *The speed of light in a vacuum used to be about 35 mph. Then Jeff Dean spent a weekend optimizing physics*

Assignment Project Exam Help

- *Jeff Dean once bit a spider, the spider got super powers and C readability*
- *Jeff Dean puts his parts on one but if he had more than two legs, you would see th h is actually $O(\log n)$*
- *Compilers don't warn Jeff Dean. Jeff Dean warns compilers*
- *The rate at which Jeff Dean produces code jumped by a factor of 40 in late 2000 when he upgraded his keyboard to USB2.0*

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Typical Big Data Problem

- Iterate over a large number of records
- Extract something of interest from each record
- Shuffle and sort intermediate results
- Aggregate intermediate results
- Generate final output

Map

Assignment Project Exam Help

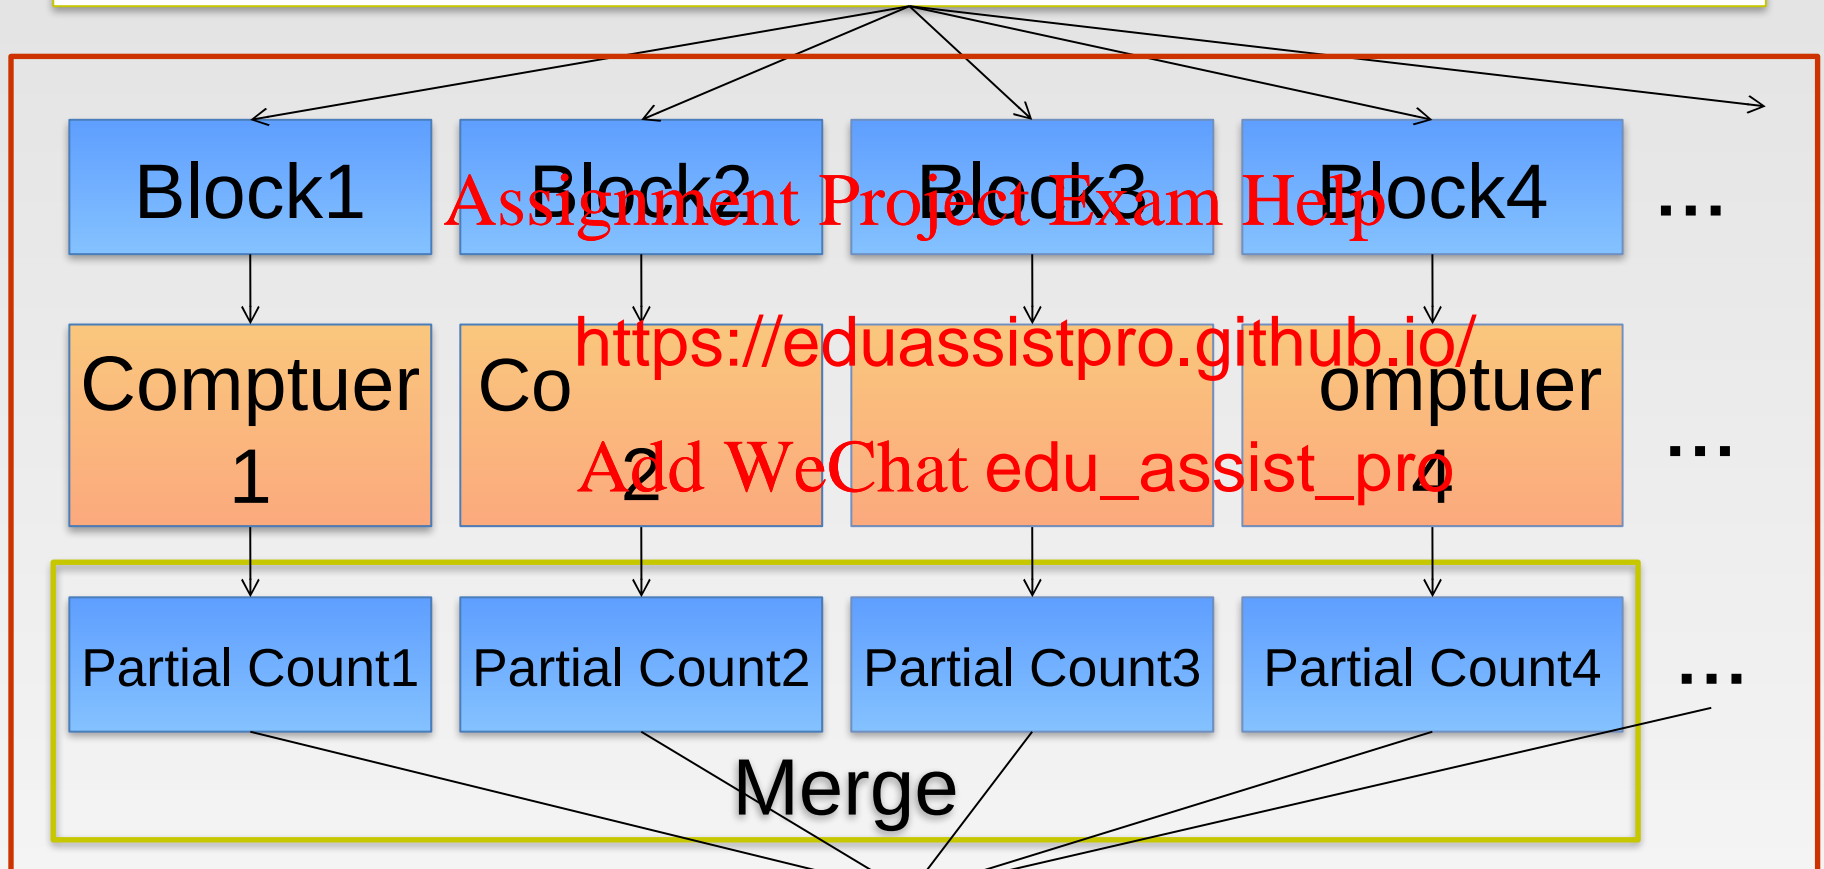
<https://eduassistpro.github.io/>

Key idea: provide abstraction
for these two operations

Add WeChat edu_assist_pro

Distributed Word Count

Huge Document



Final Result

The Idea of MapReduce

- Inspired by the map and reduce functions in functional programming
- We can view map as a transformation over a dataset
 - This transformation is specified by the function f
 - Each functional application happens in **isolation**
 - The application of f to each element of a dataset can be parallelized
- We can view reduce as an aggregation
 - The aggregation is defined by t
 - Data locality: “elements in the list are kept together”
 - If we can **group** elements of the list, also the reduce phase can proceed in parallel
- The framework coordinates the map and reduce phases:
 - Grouping intermediate results happens in parallel

Data Structures in MapReduce

- Key-value pairs are the basic data structure in MapReduce
 - Keys and values can be: integers, float, strings, raw bytes
 - They can also be arbitrary data structures
- The design of MapReduce algorithms involves
 - Imposing the structure on the input datasets
 - ▶ E.g.: for <https://eduassistpro.github.io/> keys may be URLs and values may be the HTML
 - In some algorithms, input keys uniquely identify a record (e.g., wordcount), in others they uniquely identify a record
 - Keys can be combined in complex ways to design various algorithms

Map and Reduce Functions

- Programmers specify two functions:

- **map** $(k_1, v_1) \rightarrow \text{list} [<k_2, v_2>]$

- ▶ Map transforms the input into key-value pairs to process

- **reduce** $(k_2, \text{list} [v_2]) \rightarrow [<k_3, v_3>]$

- ▶ Reduce aggregates the list of values for each key
 - ▶ All value the same reducer

- $\text{list} [<k_2, v_2>] \rightarrow \text{list} [<k_3, v_3>]$

- The MapReduce environment takes care of everything else...

- A complex program can be decomposed as a succession of Map and Reduce tasks

Everything Else?

- Handles scheduling
 - Assigns workers to map and reduce tasks
- Handles “data distribution”
 - Moves processes to data
- Handles synchronization
 - Gathers, sorts, and shuffles intermediate data
- Handles errors
 - Detects work
- Everything happens on top of a distributed file system (HDFS)
- You don’t know:
 - Where mappers and reducers run
 - When a mapper or reducer begins or finishes
 - Which input a particular mapper is processing
 - Which intermediate key a particular reducer is processing

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

A Brief View of MapReduce

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Shuffle and Sort

□ Shuffle

- Input to the Reducer is the sorted output of the mappers. In this phase the framework fetches the relevant partition of the output of all the mappers, via HTTP.

□ Sort

- The framework groups Reducer inputs by keys (since different Mappers map to the same key, this stage.

□ Hadoop framew

<https://eduassistpro.github.io/>
t step .

Add WeChat edu_assist_pro

Hadoop MapReduce Brief Data Flow

- 1. Mappers read from HDFS
- 2. Map output is partitioned by key and sent to Reducers
- 3. Reducers sort input by key
- 4. Reduce output is written to HDFS
- Intermediate results are stored on local FS of Map and Reduce workers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

“Hello World” in MapReduce

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

“Hello World” in MapReduce

- Input:
 - Key-value pairs: (docid, doc) of a file stored on the distributed filesystem
 - docid : unique identifier of a document
 - doc: is the text of the document itself
- Mapper:
 - Takes an in <https://eduassistpro.github.io/> line
 - Emits intermediate key-value p the key and the integer is the value
- The framework:
 - Guarantees all values associated with the same key (the word) are brought to the same reducer
- The reducer:
 - Receives all values associated to some keys
 - Sums the values and writes output key-value pairs: the key is the word and the value is the number of occurrences

Coordination: Master

- Master node takes care of coordination:
 - Task status: (idle, in-progress, completed)
 - Idle tasks get scheduled as workers become available
 - When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
 - Master push

<https://eduassistpro.github.io/>

- Master pings workers periodically to

Add WeChat edu_assist_pro

Dealing with Failures

- Map worker failure
 - Its task is reassigned to another map worker
 - Reduce workers are notified when task is rescheduled on another worker
- Reduce worker failure
 - Its task is reker
 - Reduce task (starting mapper tasks as well)
- Master failure
 - MapReduce task is aborted and client is notified
- Robust
 - Google's experience: lost 1600 of 1800 machines once!, but finished fine

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Where the Magic Happens

- Implicit between the map and reduce phases is a **parallel “group by”** operation on intermediate keys
 - Intermediate data arrive at each reducer in order, sorted by the key
 - No ordering is guaranteed across reducers
- Output keys from <https://eduassistpro.github.io/> DFS
 - The output may consist of r distributed files where r is the number of reducers
 - Such output may be the input to a subsequent MapReduce phase
- Intermediate keys (used in shuffle and sort) are transient:
 - They are not stored on the distributed filesystem
 - They are “spilled” to the local disk of each machine in the cluster

Assignment Project Exam Help

Write Your Program in Java? <https://eduassistpro.github.io>

Add WeChat edu_assist_pro

MapReduce Program

- A MapReduce program consists of the following 3 parts:
 - Driver → main (would trigger the map and reduce methods)
 - Mapper
 - Reducer
 - It is better to include the map reduce and main methods in 3 different cla

<https://eduassistpro.github.io/>

- Check detailed information of all cla

<https://hadoop.apache.org/docs/r2.2.4/hadoop-mapreduce-client-jobmodel.html>

Mapper

```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
```

Assignment Project Exam Help

```
    public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

Mapper Explanation

- Maps input key/value pairs to a set of intermediate key/value pairs.

//Map class header

```
public static class TokenizerMapper
```

```
    extends Mapper<Object, Text, Text, IntWritable>{
```

□ Class Mapper<KEYIN,VALUEIN,KEYOUT,VALUEOUT>

- ▶ KEYIN,V oc)

- ▶ KEYOUT rd,1)

Add WeChat edu_assist_pro

// IntWritable: A serializable and comparable integer

```
private final static IntWritable one = new IntWritable(1);
```

//Text: stores text using standard UTF8 encoding. It provides methods to serialize, deserialize, and compare texts at byte level

```
private Text word = new Text();
```

//hadoop supported data types for the key/value pairs, in package org.apache.hadoop

What is Writable?

- Hadoop defines its own “box” classes for strings (Text), integers (IntWritable), etc.

- All values must implement interface Writable

Assignment Project Exam Help

- All keys must implement WritableComparable

<https://eduassistpro.github.io/>

- Writable is a serializable object which implements a simple, efficient, serialization protocol

Add WeChat edu_assist_pro

Mapper Explanation (Cont')

//Map method header

public void map(Object key, Text value, Context context) throws
IOException, InterruptedException

- Object key/Text value: Data type of the input Key and Value to the mapper
- Context: An inner class of Mapper, used to store the context of a running task
Mapper or the output by either the inputs or the output of the job
- Exceptions: IOException, InterruptedException
- This function is called once for each key/value pair in the input split. Your application should override this to do your job.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Mapper Explanation (Cont')

```
//Use a string tokenizer to split the document into words
StringTokenizer itr = new StringTokenizer(value.toString());
//Iterate through each word and a form key value pairs
while (itr.hasMoreTokens()) {
    //Assign each word to a Text object of String type, it a Text 'word'
    word.set(itr.nextToken());
    //Form key value pair using context
    context.write(word, one);
}
```

- Map function produces Map.Context object
 - Map.context() takes (k, v) elements
- Any (WritableComparable, Writable) can be used

Reducer

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException{
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Reducer Explanation

//Reduce Header similar to the one in map with different key/value data type

```
public static class IntSumReducer
```

```
    extends Reducer<Text, IntWritable, Text, IntWritable>
```

//data from map will be <"word",{1,1,..}>, so we get it with an Iterator and thus we can go through the sets of values

```
public void reduce(Text key, Iterable<IntWritable> values,
Context context) throws IOException, InterruptedException {
```

//Initaize a variable 'sum' <https://eduassistpro.github.io/>

```
    int sum = 0;
```

//Iterate through all the values with respect to key and add them up

```
    for (IntWritable val : values) {
```

```
        sum += val.get();
```

```
    }
```

// Form the final key/value pairs results for each word using context

```
    result.set(sum);
```

```
    context.write(key, result);
```

Main (Driver)

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizerMapper.class);  
    job.setReducerClass(Reducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Main(The Driver)

- Given the Mapper and Reducer code, the short main() starts the MapReduction running
- The Hadoop system picks up a bunch of values from the command line on its own
- Then the main() also specifies a few key parameters of the problem in the Job object
- Job is the primary class that defines a map-reduce job to the Hadoop framework. It specifies what Map and Reduce classes to use and the format of the input and output files)
- Other parameters, i.e. the number of reducers, are optional and the system will determine good values for them if not specified
- Then the framework tries to faithfully execute the job as-is described by Job

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Main Explanation

//Creating a Configuration object and a Job object, assigning a job name for identification purposes

```
Configuration conf = new Configuration();
```

```
Job job = Job.getInstance(conf, "word count");
```

- Job Class: It allows the user to configure the job, submit it, control its execution, and query the state. Normally the user creates the application, job via `Job` and then submits the

//Setting the job's jar file by finding the provi

```
job.setJarByClass(WordCount.class);
```

//Providing the mapper and reducer class names

```
job.setMapperClass(TokenizerMapper.class);
```

```
job.setReducerClass(IntSumReducer.class);
```

//Setting configuration object with the Data Type of output Key and Value for map and reduce

```
job.setOutputKeyClass(Text.class);
```

```
job.setOutputValueClass(IntWritable.class);
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Main Explanation (Cont')

//The hdfs input and output directory to be fetched from the command line

```
FileInputFormat.addInputPath(job, new Path(args[0]));
```

```
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

//Submit the job to the cluster and wait for it to finish.

```
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Make It Running !

- Configure environment variables

```
export JAVA_HOME=...
```

```
export PATH=${JAVA_HOME}/bin:${PATH}
```

```
export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
```

- Compile WordCount.java and create a jar.

```
$ hadoop com.sun.tools.j
```

```
$ jar cf wc.jar WordCount
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Put files to HDFS

```
$ hdfs dfs -put YOURFILES input
```

Add WeChat edu_assist_pro

- Run the application

```
$ hadoop jar wc.jar WordCount input output
```

- Check the results

```
$ hdfs dfs -cat output/*
```

Make It Running !

- Given two files:
 - file1: Hello World Bye World
 - file2: Hello Hadoop Goodbye Hadoop
- The first map emits:
 - < Hello, 1> < World, 1> < Bye, 1> < World, 1>
- The second map emits:
 - < Hello, 1> < Hadoop, 1>
- The output of the job is:
 - < Bye, 1> < Goodbye, 1> < Hadoop, 2> < World, 2>

Mappers and Reducers

- Need to handle more data? Just add more Mappers/Reducers!
- No need to handle multithreaded code ☺
 - Mappers and Reducers are typically single threaded and deterministic
 - ▶ Determinism allows for restarting of failed jobs
 - Mappers/Reducers don't know anything about each other
 - ▶ In Hadoop

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Combiners

- Often a Map task will produce many pairs of the form $(k, v_1), (k, v_2), \dots$ for the same key k
 - E.g., popular words in the word count example
- Combiners are a general mechanism to reduce the amount of intermediate data, thus saving network time
 - They could be thought of as “mini-reducers”
- Warning!
 - The use of combiners must be **https://eduassistpro.github.io/** ly
 - ▶ Optional in Hadoop. the combiner algorithm **cannot depend on** computation (or even execution) of the combiners
 - ▶ A combiner operates on each map output key. It must have the same output key-value types as the Mapper class.
 - ▶ A combiner can produce summary information from a large dataset because it replaces the original Map output
 - Works only if reduce function is commutative and associative
 - ▶ In general, reducer and combiner **are not interchangeable**

Combiners in WordCount

- Combiner combines the values of all keys of a single mapper node (single machine):

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

- Much less data needs to be copied and shuffled!
- If combiners take advantage of all opportunities for local aggregation we have at most $m \times V$ intermediate key-value pairs
 - m : number of mappers
 - V : number of unique terms in the collection
- Note: not all mappers will see all terms

Combiners in WordCount

- In WordCount.java, you only need to add the follow line to Main:

```
job.setCombinerClass(IntSumReducer.class);
```

- This is because in this example, Reducer and Combiner do the same thing

- Note: Most cases this is not true!

- You need to

- Given two files: <https://eduassistpro.github.io/>

- file1: Hello World Bye World

- file2: Hello Hadoop Goodbye H

- The first map emits:

- < Hello, 1> < World, 2> < Bye, 1>

- The second map emits:

- < Hello, 1> < Hadoop, 2> < Goodbye, 1>

Partitioner

- Partitioner controls the partitioning of the keys of the intermediate map-outputs.
 - The key (or a subset of the key) is used to derive the partition, typically by a *hash function*.
 - The total number of partitions is the same as the number of reduce tasks for the job.
 - ▶ This controls the intermediate key (and hence the number of partitions) using the *hash function*.
- System uses HashPartitioner by default
 - $\text{hash}(\text{key}) \bmod R$
- Sometimes useful to override the hash function:
 - E.g., ***hash(hostname(URL)) mod R*** ensures URLs from a host end up in the same output file
 - ▶ <https://www.unsw.edu.au/faculties> and <https://www.unsw.edu.au/about-us> will be stored in one file
- Job sets Partitioner implementation (in Main)

MapReduce: Recap

□ Programmers must specify:

□ $\text{map}(k_1, v_1) \rightarrow [(k_2, v_2)]$

□ $\text{reduce}(k_2, [v_2]) \rightarrow [k_3, v_3]$

□ All values with the same key are reduced together

□ Optionally, also:

□ $\text{combine}(k_2,$

▶ Mini-red

▶ Used as an optimization to r

□ $\text{partition}(k_2, \text{number of partition}$

▶ Often a simple hash of the key, e.g., $\text{hash}(k_2) \bmod n$

▶ Divides up key space for parallel reduce operations

□ The execution framework handles everything else...

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

e map phase

traffic

k_2

MapReduce: Recap (Cont')

- ▢ Divides input into fixed-size pieces, ***input splits***
 - ▢ Hadoop creates one map task for each split
 - ▢ Map task runs the user-defined map function for each *record* in the split
- ▢ Size of splits
 - ▢ Small size is r machine will be able to process
<https://eduassistpro.github.io/>
 - ▢ But if splits are too small, the ov aging the splits
dominate the total execution time
 - ▢ For most jobs, a good split size size of a HDFS block, 64MB(default)
- ▢ Data locality optimization
 - ▢ Run the map task on a node where the input data resides in HDFS
 - ▢ This is the reason why the split size is the same as the block size

MapReduce: Recap (Cont')

- Map tasks write their output to local disk (not to HDFS)
 - Map output is intermediate output
 - Once the job is complete the map output can be thrown away
 - So storing it in HDFS with replication would be overkill
 - If the node of map task fails, Hadoop will automatically rerun the map task on
- Reduce tasks d <https://eduassistpro.github.io/> locally
 - Input to a single reduce task is tput from all mappers Add WeChat edu_assist_pro
 - Output of the reduce is stored in HDFS for reliability
- The number of reduce tasks is not governed by the size of the input, but is specified independently

MapReduce: Recap (Cont')

- When there are multiple reducers, the map tasks partition their output:
 - One partition for each reduce task
 - The records for every key are all in a single partition
 - Partitioning can be controlled by a user-defined partitioning function

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

More Detailed MapReduce Dataflow

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

MapReduce: Recap

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Another Example : Analysis of Weather Dataset

- Data from NCDC(National Climatic Data Center)
 - A large volume of log data collected by weather sensors: e.g. temperature
- Data format
 - Line-oriented ASCII format
 - Each record has many elements
 - We focus on the temperature element
 - Data files are organized
 - There is a directory for each weather station with its readings for the year
- Query
 - What's the highest recorded global temperature for each year in the dataset?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Year	Temperature
00670119909999991950051507004...	9999999N9+00001+9999999999
9...	
00430119909999991950051512004...	9999999N9+00221+9999999999
9...	
00430119909999991950051518004...	9999999N9-
001111+999999999999...	
00430126509999991949032412004...	0500001N9+01111+9999999999
9...	
00430126509999991949032418004...	0500001N9+00781+9999999999

Contents of data files

```
% ls raw/1990 | head
010010-99999-1990.gz
010014-99999-1990.gz
010015-99999-1990.gz
010016-99999-1990.gz
010017-99999-1990.gz
010030-99999-1990.gz
010040-99999-1990.gz
010080-99999-1990.gz
010100-99999-1990.gz
010150-99999-1990.gz
```

List of data files

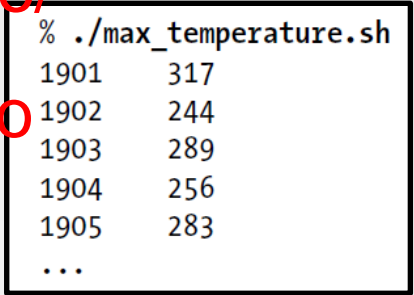
Analyzing the Data with Unix Tools

- To provide a performance baseline
- Use *awk* for processing line-oriented data
- Complete run for the century took **42 minutes** on a single EC2 High-CPU Extra Large Instance

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



A terminal window showing the execution of a script. The prompt is '% ./max_temperature.sh'. The output consists of five lines of data, each with a year and a temperature value, followed by an ellipsis indicating more data.

1901	317
1902	244
1903	289
1904	256
1905	283
...	

How Can We Parallelize This Work?

- To speed up the processing, we need to run parts of the program in **parallel**
- **Challenges?**
 - Divide the work into even distribution is not easy
 - ▶ File size for different years varies
 - Combining the results of each chunk
 - ▶ Get the right nature for each chunk
 - We are still limited by the processing power of a single machine
 - ▶ Some datasets grow beyond a single machine
- To use **multiple machines**, we need to consider a variety of complex problems
 - Coordination: Who runs the overall job?
 - Reliability: How do we deal with failed processes?
- **Hadoop** can take care of these issues

MapReduce Design

□ We need to answer these questions:

□ What are the map input key and value types?

□ What does the mapper do?

□ What are the map output key and value types?

□ Can we use a combiner?

□ Is a partition

□ What does t <https://eduassistpro.github.io/>

□ What are the reduce output key s?

Add WeChat edu_assist_pro

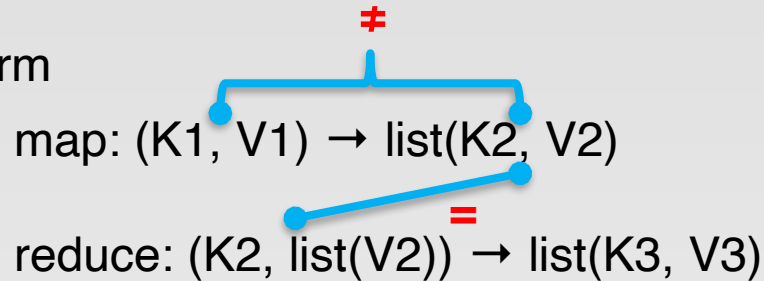
□ And: What are the file formats?

□ For now we are using text files

□ We may use binary files

MapReduce Types

- General form



- Combine function

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- The same form as the reduce function, its output types
- Output type is the same as Map
- The combine and reduce functions may be the same

- Partition function

partition: $(K2, V2) \rightarrow \text{integer}$

- Input intermediate key and value types
- Returns the partition index

MapReduce Design

- Identify the input and output of the problem
 - Text input format of the dataset files (input of mapper)
 - ▶ Key: offset of the line (unnecessary)
 - ▶ Value: each line of the files (string)
 - Output (output of reducer)
 - ▶ Key: yea
 - ▶ Value: m
- Decide the MapReduce data types
 - Hadoop provides its own set of
 - ▶ optimized for network serialization
 - ▶ org.apache.hadoop.io package
 - In WordCount, we have used Text and IntWritable
 - Key must implement interface WritableComparable
 - Value must implement interface Writable

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Writable Wrappers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Writable Class Hierarchy

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

What does the Mapper Do?

- Pull out the year and the temperature
 - Indeed in this example, the map phase is simply data preparation phase
 - Drop bad records(filtering)

Assignment Project Exam Help

Input File

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

of Map Function (key, value)

Input of Map Function (key, value)

```
(0, 0067011990999991950051507004...9999999N9+00001+9999999999...)  
(106, 0043011990999991950051512004...9999999N9+00221+9999999999...)  
(212, 0043011990999991950051518004...9999999N9-00111+9999999999...)  
(318, 0043012650999991949032412004...0500001N9+01111+9999999999...)  
(424, 0043012650999991949032418004...0500001N9+00781+9999999999...)
```

Map



```
(1950, 0)  
(1950, 22)  
(1950, -11)  
(1949, 111)  
(1949, 78)
```

Map Input and Output

□ Input

□ Key: offset of the line (unnecessary)

- ▶ The dataset is quite large and contains a huge number of lines
- ▶ LongWritable

□ Value: each line of the files (string)

- ▶ Text

□ Output

□ Key: year

- ▶ Both string or integer format
- ▶ Text/IntWritable

□ Value: temperature

- ▶ Integer is already enough to store it
- ▶ IntWritable

□ Combiner and Partitioner?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

What does the Reducer Do?

- ❑ Reducer input
 - ❑ (year, [temperature1, temperature2, temperature3, ...])
- ❑ Scan all values received for the key, and find out the maximum one

Assignment Project Exam Help

- ❑ Reducer output
 - ❑ Key: year <https://eduassistpro.github.io/>
 - ▶ String/IntWritable
 - ❑ Value: maximum temperature
 - ▶ IntWritable

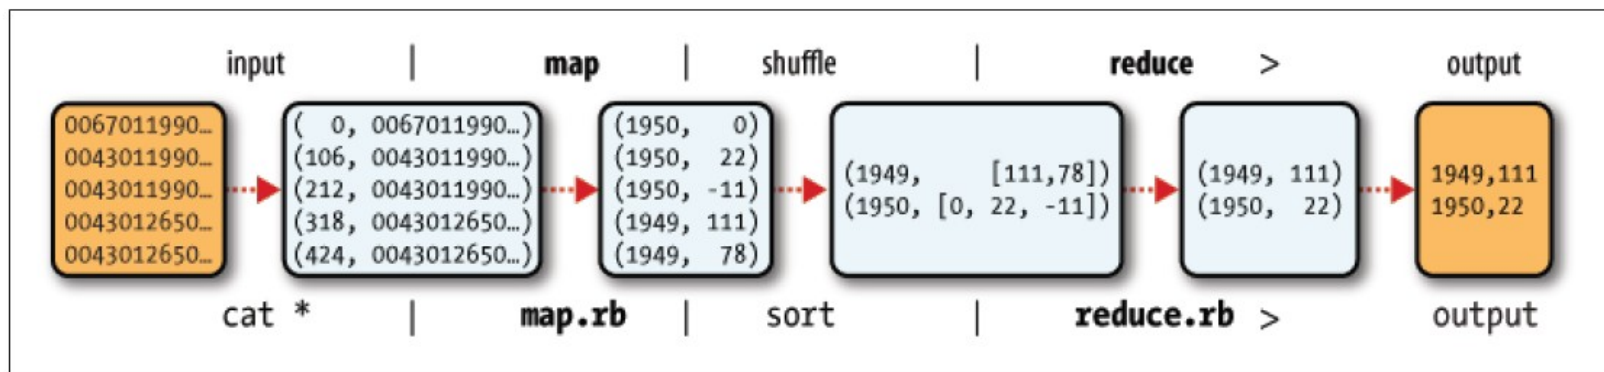
Add WeChat edu_assist_pro

MapReduce Design of NCDC Example

- The output from the map function is processed by MapReduce framework
 - Sorts and groups the key-value pairs by key



- Reduce function iterates over the maximum value



Java Implementation of the Example

```
public class MaxTemperatureMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private static final int MISSING = 9999;  
  
    @Override  
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {  
        String line = value.toString();  
        String year = line.substring(15, 19);  
        int airTemperature;  
        if (line.charAt(87) == '+') {  
            airTemperature = Integer.parseInt(line.substring(87, 92));  
        } else {  
            airTemperature = Integer.parseInt(line.substring(87, 92));  
        }  
        String quality = line.substring(92, 93);  
        if (airTemperature != MISSING && quality.matches("[01459]")) {  
            context.write(new Text(year), new IntWritable(airTemperature));  
        }  
    }  
}
```

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Java Implementation of the Example

```
public class MaxTemperatureReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int maxValue = Integer.MIN_VALUE;
        for (IntWritable value : values) {
            if (value.get() > maxValue)
                maxValue = value.get();
        }
        context.write(key, new IntWritable(maxValue));
    }
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Java Implementation of the Example

```
public class MaxTemperatureWithCombiner {  
    //specify the usage of the job  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.err.println("Usage: MaxTemperatureWithCombiner <input path> " + "<output path>");  
            System.exit(-1);  
        }  
        //Construct a job object to configure, control and run the job  
        Job job = new Job();  
        job.setJarByClass(MaxTemp  
        job.setJobName("Max temper  
        //Specify input and output path  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        //Specify map and reduce classes, also a combiner  
        job.setMapperClass(MaxTemperatureMapper.class);  
        job.setCombinerClass(MaxTemperatureReducer.class);  
        job.setReducerClass(MaxTemperatureReducer.class);  
        //Specify output type  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        //submit the job and wait for completion  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Codes can be found here:

<http://hadoopbook.com/code.html>

References

- Chapter 2, Hadoop The Definitive Guide

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro