

COMP9313: Big Data Management

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Assignment Project Exam Help

Ch <https://eduassistpro.github.io/> **cell**

Add WeChat edu_assist_pro

Overview of Previous Lecture

- Motivation of MapReduce
- Data Structures in MapReduce: (key, value) pairs
- Map and Reduce Functions
- Hadoop MapReduce Programming
 - Mapper
 - Reducer
 - Combiner
 - Partitioner
 - Driver

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Combiner Function

- To minimize the data transferred between map and reduce tasks
- Combiner function is run on the map output
- Both input and output data types must be consistent with the output of mapper (or input of reducer)
- But Hadoop do not guarantee how many times it will call combiner function for a particular map output record
 - It is just opti
 - The number of calling (even zero) does not affect the output of Reducers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

$$\max(0, 20, 10, 25, 15) = \max(\max(0, 20, 10), \max(25, 15)) = \max(20, 25) = 25$$

- Applicable on problems that are commutative and associative
 - Commutative: $\max(a, b) = \max(b, a)$
 - Associative: $\max(\max(a, b), c) = \max(a, \max(b, c))$

Assignment Project Exam Help

MapReduce <https://eduassistpro.github.io/> **Design Patterns**

Add WeChat edu_assist_pro

Assignment Project Exam Help

Design Pa <https://eduassistpro.github.io/> **n Combining**

Add WeChat edu_assist_pro

Importance of Local Aggregation

- Ideal scaling characteristics:
 - Twice the data, twice the running time
 - Twice the resources, half the running time
- Why can't we achieve this?
 - Data synchronization requires communication
 - Communication
- Thus... avoid communication!
 - Reduce intermediate data via local aggregation
 - Combiners can help

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

WordCount Baseline

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

What's the impact of combiners?

Word Count: Version 1

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Are combinators still needed?

Word Count: Version 2

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Design Pattern for Local Aggregation

- “In-mapper combining”
 - Fold the functionality of the combiner into the mapper by preserving state across multiple map calls
- Advantages
 - Speed
 - Why is this f <https://eduassistpro.github.io/>
- Disadvantages
 - Explicit memory management required
 - Potential for order-dependent bugs

Assignment Project Exam Help

Add WeChat edu_assist_pro

Combiner Design

- Combiners and reducers share same method signature
 - Sometimes, reducers can serve as combiners
 - Often, not...

- Remember: combiners are optional optimizations

- Should not

- May be run <https://eduassistpro.github.io/>

- Example: find average of all integer with the same key

Computing the Mean: Version 1

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Why can't we use reducer as combiner?

$\text{Mean}(1, 2, 3, 4, 5) \neq \text{Mean}(\text{Mean}(1, 2), \text{Mean}(3, 4, 5))$

Computing the Mean: Version 2

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Why doesn't this work? Combiners must have the same input and output type, consistent with the input of reducers (output of mappers)

Computing the Mean: Version 3

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Fixed?

Check the correctness by removing the combiner

Computing the Mean: Version 4

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Assignment Project Exam Help

How to Implement In-mapper Combiner

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Lifecycle of Mapper/Reducer

- Lifecycle: setup -> map -> cleanup
 - setup(): called once at the beginning of the task
 - map(): do the map
 - cleanup(): called once at the end of the task.
 - We do not invoke these functions
 - In-mapper Com
 - Use setup() <https://eduassistpro.github.io/> data structure
 - Use cleanup() to emit the final
- Add WeChat edu_assist_pro

Word Count: Version 2

setup()

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

cleanup()

Assignment Project Exam Help

Design <https://eduassistpro.github.io/> **Stripe**

Add WeChat edu_assist_pro

Term Co-occurrence Computation

- Term co-occurrence matrix for a text collection
 - $M = N \times N$ matrix (N = vocabulary size)
 - M_{ij} : number of times i and j co-occur in some context (for concreteness, let's say context = sentence)
 - specific instance of a large counting problem
 - ▶ A large e
 - ▶ A large n (section itself)
 - ▶ Goal: keep track of interesti out the events
- Basic approach
 - Mappers generate partial counts
 - Reducers aggregate partial counts
- How do we aggregate partial counts efficiently?

First Try: “Pairs”

- Each mapper takes a sentence
 - Generate all co-occurring term pairs
 - For all pairs, emit (a, b) → count
- Reducers sum up counts associated with these pairs
- Use combiners

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

“Pairs” Analysis

□ Advantages

- Easy to implement, easy to understand

□ Disadvantages

- Lots of pairs to sort and shuffle around (upper bound?)
- Not many o rk

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Another Try: “Stripes”

- Idea: group together pairs into an associative array

(a, b) \rightarrow 1

(a, c) \rightarrow 2

(a, d) \rightarrow 5

(a, e) \rightarrow 3

(a, f) \rightarrow 2

$a \rightarrow \{ b: 1, c: 2, d: 5, e: 3, f: 2 \}$

Assignment Project Exam Help

- Each mapper takes <https://eduassistpro.github.io/>
 - Generate all co-occurring terms
 - For each term, emit $a \rightarrow \{ b: \text{count}_b, c: \text{count}_c, d: \text{count}_d \dots \}$
- Reducers perform element-wise sum of associative arrays

$$\begin{array}{rcl} & a \rightarrow \{ b: 1, & d: 5, e: 3 \} \\ + & a \rightarrow \{ b: 1, c: 2, d: 2, & f: 2 \} \\ \hline & a \rightarrow \{ b: 2, c: 2, d: 7, e: 3, f: 2 \} \end{array}$$

Key: cleverly-constructed data structure
brings together partial results

Stripes: Pseudo-Code

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

“Stripes” Analysis

□ Advantages

- Far less sorting and shuffling of key-value pairs
- Can make better use of combiners

□ Disadvantages

- More difficult
- Underlying
- Fundamental limitation in terms of space

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Compare “Pairs” and “Stripes”

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Cluster size: 38 cores

Data Source: Associated Press Worldstream (APW) of the English Gigaword Corpus (v3), which contains 2.27 million documents (1.8 GB compressed, 5.7 GB uncompressed)

Pairs vs. Stripes

- The pairs approach
 - Keep track of each team co-occurrence separately
 - Generates a large number of key-value pairs (also intermediate)
 - The benefit from combiners is limited, as it is less likely for a mapper to process multiple occurrences of a word
- The stripe approach
 - Keep track of the same term
 - Generates fewer and shorter intermediate keys
 - The framework has less sorting
 - Greatly benefits from combiners, as the key space is the vocabulary
 - More efficient, but may suffer from memory problem
- These two design patterns are broadly useful and frequently observed in a variety of applications
 - Text processing, data mining, and bioinformatics

Assignment Project Exam Help **How to Implement "Pairs" and "Stripes"**

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Serialization

- Process of turning structured objects into a byte stream for transmission over a network or for writing to persistent storage

- Deserialization is the reverse process of serialization

Assignment Project Exam Help

- Requirements

- Compact <https://eduassistpro.github.io/>

- ▶ To make efficient use of sto

- Fast [Add WeChat edu_assist_pro](#)

- ▶ The overhead in reading and writing of data is minimal

- Extensible

- ▶ We can transparently read data written in an older format

- Interoperable

- ▶ We can read or write persistent data using different language

Writable Interface

- ❑ Hadoop defines its own “box” classes for strings (Text), integers (IntWritable), etc.
- ❑ Writable is a serializable object which implements a simple, efficient, serialization protocol

```
public interface Writable {  
    void wri  
    void rea  
}
```

- ❑ All values must implement interface
- ❑ All keys must implement interface WritableComparable
- ❑ context.write(WritableComparable, Writable)
 - ❑ You cannot use java primitives here!!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Writable Wrappers for Java Primitives

- There are **Writable** wrappers for all the Java primitive types except short and char (both of which can be stored in an **IntWritable**)
- **get()** for retrieving and **set()** for storing the wrapped value
- Variable-length formats
 - If a value is between -128 and 127, use only a single byte
 - Otherwise, use two bytes if the value is positive or negative

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Writable Examples

□ Text

- Writable for UTF-8 sequences
- Can be thought of as the Writable equivalent of `java.lang.String`
- Maximum size is 2GB
- Use standard UTF-8
- Text is mutable operations, except
NullWritable <https://eduassistpro.github.io/>
 - ▶ Different from `java.lang.String`
 - ▶ You can reuse a Text instance of the `set()` method

□ NullWritable

- Zero-length serialization
- Used as a placeholder
- A key or a value can be declared as a **NullWritable** when you don't need to use that position

Stripes Implementation

□ A stripe key-value pair $a \rightarrow \{b: 1, c: 2, d: 5, e: 3, f: 2\}$:

□ Key: the term a

□ Value: the stripe $\{b: 1, c: 2, d: 5, e: 3, f: 2\}$

▶ In Java, easy, use map (hashmap)

▶ How to represent this stripe in MapReduce?

□ MapWritable: the `put` method

□ `put(Object key, Writable value)`

□ `get(Object key)`

□ `containsKey(Object key)`

□ `containsValue(Writable value)`

□ `entrySet()` , returns

`Set<Map.Entry<Writable,Writable>>`, used for iteration

□ More details please refer to

<https://hadoop.apache.org/docs/r2.7.2/api/org/apache/hadoop/io/MapWritable.html>

Pairs Implementation

□ Key-value pair (a, b) → count

□ Value: count

□ Key: (a, b)

▶ In Java, easy, implement a pair class

▶ *How to store the key in Map Reduce?*

Assignment Project Exam Help

□ You must custom <https://eduassistpro.github.io/> Implement interface WritableComparable!

Add WeChat edu_assist_pro

□ First start from a easier task: when the value is a pair, which must implement interface Writable

Multiple Output Values

- If we are to output multiple values for each key
 - E.g., a pair of String objects, or a pair of int
- How do we do that?
- WordCount output a single number as the value
- Remember, our object containing the values needs to implement the Writable interface
- We could use `T`
 - Value is a string of comma sep
 - Have to convert the values to string
 - Have to parse the string on input (not hard) to get the values

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

Implement a Custom Writable

- Suppose we wanted to implement a custom class containing a pair of integers. Call it IntPair.
- How would we implement this class?
 - Needs to implement the Writable interface
 - Instance variables to hold the values
 - Construct fu
 - A method to <https://eduassistpro.github.io/>
 - A method to get the values (two
 - write() method : serialize the les (two integers) objects in turn to the output stream
 - readFields() method: deserialize the member variables (two integers) in turn from the input stream
 - As in Java: hashCode(), equals(), toString()

Implement a Custom Writable

- Implement the Writable interface

```
public class IntPair implements Writable {
```

- Instance variables to hold the values

```
    private int first, second;
```

- Construct functions

```
    public IntPair() {  
    }  
  
    public IntPair(int first, int second) {  
        set(first, second);  
    }  
}
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- set() method

```
    public void set(int left, int right) {  
        first = left;  
        second = right;  
    }  
}
```

Implement a Custom Writable

□ get() method

```
public int getFirst() {  
    return first;  
}  
public int getSecond() {  
    return second;  
}
```

Assignment Project Exam Help

□ write() method

```
public void write(D  
    out.writeInt(first);  
    out.writeInt(second);  
}
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

□ Write the two integers to the output stream in turn

□ readFields() method

```
public void readFields(DataInput in) throws IOException {  
    first = in.readInt();  
    second = in.readInt();  
}
```

□ Read the two integers from the input stream in turn

Complex Key

- If the key is not a single value
 - E.g., a pair of String objects, or a pair of int
- How do we do that?
- The co-occurrence matrix problem, a pair of terms as the key
- Our object containing the values needs to implement the WritableComparable
 - Why WritableComparable? <https://eduassistpro.github.io/>
- We could use Text again
 - Value is a string of comma sep
 - Have to convert the values to strings, build the full string
 - Have to parse the string on input (not hard) to get the values
 - Objects are compared according to the full string!!

Implement a Custom WritableComparable

- Suppose we wanted to implement a custom class containing a pair of String objects. Call it StringPair.
- How would we implement this class?
 - Needs to implement the WritableComparable interface
 - Instance variables to hold the values
 - Construct fu
 - A method to <https://eduassistpro.github.io/>
 - A method to get the values (two)
 - write() method : serialize the les (i.e., two String) objects in turn to the output stream
 - readFields() method: deserialize the member variables (i.e., two String) in turn from the input stream
 - As in Java: hashCode(), equals(), toString()
 - compareTo() method: specify how to compare two objects of the self-defined class

Implement a Custom WritableComparable

- implement the Writable interface

```
public class StringPair implements WritableComparable<StringPair> {
```

- Instance variables to hold the values

```
    private String first, second;
```

- Construct functions

```
    public StringPair()  
    {
```

```
        public StringPair(String first, String second)  
        {  
            set(first, second);  
        }  
    }
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- set() method

```
    public void set(String left, String right) {  
        first = left;  
        second = right;  
    }
```

Implement a Custom WritableComparable

□ get() method

```
public String getFirst() {  
    return first;  
}  
public String getSecond() {  
    return second;  
}
```

Assignment Project Exam Help

□ write() method

```
public void write(D  
    String[] strings = new String[] { first, sec  
WritableUtils.writeStringArray(out, strings  
}
```

<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](#)

□ Utilize WritableUtils.

□ readFields() method

```
public void readFields(DataInput in) throws IOException {  
    String[] strings = WritableUtils.readStringArray(in);  
    first = strings[0];  
    second = strings[1];  
}
```

Implement a Custom WritableComparable

□ compareTo() method:

```
public int compareTo(StringPair o) {  
    int cmp = compare(first, o.getFirst());  
    if(cmp != 0){  
        return cmp;  
    }  
    return compare(second, o.getSecond());  
}
```

```
private int compare(String s1, String s2) {  
    if (s1 == null && s2 == null) {  
        return 0;  
    } else if (s1 != null && s2 == null) {  
        return 1;  
    } else if (s1 == null && s2 != null) {  
        return -1;  
    } else {  
        return s1.compareTo(s2);  
    }  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Implement a Custom WritableComparable

- You can also make the member variables as Writable objects
- Instance variables to hold the values

```
private Text first, second;
```

- Construct functions

```
public StringPair() {  
    set(new Text(), new Text());  
}
```

```
public StringPair(Text first, Text second) {  
    set(first, second);  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- set() method

```
public void set(Text left, Text right) {  
    first = left;  
    second = right;  
}
```

Implement a Custom WritableComparable

□ get() method

```
public Text getFirst() {  
    return first;  
}  
public Text getSecond() {  
    return second;  
}
```

Assignment Project Exam Help

□ write() method

```
public void write(D  
    first.write(out);  
    second.write(out);  
}
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

□ Delegated to Text

□ readFields() method

```
public void readFields(DataInput in) throws IOException {  
    first.readFields(in);  
    second.readFields(in);  
}
```

□ Delegated to Text

Implement a Custom WritableComparable

- In some cases such as secondary sort, we also need to override the hashCode() method.
- Because we need to make sure that all key-value pairs associated with the first part of the key are sent to the same reducer!

```
public int hashCode()  
    return first.hashCode();  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- By doing this, a partitioner will only use the hashCode() of the first part.
- You can also write a partitioner to do this job

Assignment Project Exam Help

Design <https://eduassistpro.github.io/> **Inversion**

Add WeChat edu_assist_pro

Computing Relative Frequencies

- “Relative” Co-occurrence matrix construction
 - Similar problem as before, same matrix
 - Instead of absolute counts, we take into consideration the fact that some words appear more frequently than others
 - ▶ Word w_i may co-occur frequently with word w_j simply because one of the two is very common
 - We need to compute relative frequencies $f(w_j|w_i)$
 - ▶ What proportion of the time w_j occurs in the context of w_i ?
- Formally, we compute:

$$f(w_j|w_i) = \frac{N(w_i, w_j)}{\sum_{w'} N(w_i, w')}$$

- $N(\cdot, \cdot)$ is the number of times a co-occurring word pair is observed
- The denominator is called the marginal

$f(w_j|w_i)$: “Stripes”

- In the reducer, the counts of all words that co-occur with the conditioning variable (w_i) are available in the associative array

- Hence, the sum of all those counts gives the marginal

Assignment Project Exam Help

- Then we divide t and we're done

<https://eduassistpro.github.io/>

$a \rightarrow \{b_1:3, b_2:12, b_3:7,$

Add WeChat edu_assist_pro

$$f(b_1|a) = 3 / (3 + 12 + 7 + 1 + \dots)$$

- Problems?

- Memory

$f(w_j|w_i)$: “Pairs”

- The reducer receives the pair (w_i, w_j) and the count
- From this information alone it is not possible to compute $f(w_j|w_i)$
 - Computing relative frequencies requires marginal counts
 - But the marginal cannot be computed until you see all counts

Assignment Project Exam Help

$((a, b_1), \{1, 1,$ <https://eduassistpro.github.io/>

[Add WeChat edu_assist_pro](#)
No way to compute $f(b_1|a)$ because a is unknown

$f(w_j | w_i) : \text{“Pairs”}$

- Solution 1: Fortunately, as for the mapper, also the reducer can preserve state across multiple keys
 - We can buffer in memory all the words that co-occur with w_i and their counts
 - This is basically building the associative array in the stripes method

Assignment Project Exam Help

<https://eduassistpro.github.io/>

$a \rightarrow \{b_1:3, b_2:12, b_3:$

Add WeChat edu_assist_pro

is now buffered in the reducer side

□ Problems?

$f(w_j | w_i) : \text{“Pairs”}$

If reducers receive pairs not sorted

$((a, b_1), \{1, 1, 1, \dots\})$

$((c, d_1), \{1, 1, 1, \dots\})$

$((a, b_2), \{1, 1, 1, \dots\})$

...

<https://eduassistpro.github.io/>
W original?

Add WeChat edu_assist_pro

- We must define the sort order of the pairs
 - In this way, the keys are first sorted by the left word, and then by the right word (in the pair)
 - Hence, we can detect if all pairs associated with the word we are conditioning on (w_i) have been seen
 - At this point, we can use the in-memory buffer, compute the relative frequencies and emit

$f(w_j | w_i)$: “Pairs”

$((a, b_1), \{1, 1, 1, \dots\})$ and $((a, b_2), \{1, 1, 1, \dots\})$ may be assigned to different reducers!

Default partitioner computed based on the whole key.

Assignment Project Exam Help

- We must define
 - The default <https://eduassistpro.github.io/> value of the intermediate key, modulo the n reducers
 - For a complex key, the raw byte array is used to compute the hash value
 - ▶ Hence, there is no guarantee that the pair (dog, aardvark) and (dog, zebra) are sent to the same reducer
 - What we want is that all pairs with the same left word are sent to the same reducer
- Still suffer from the memory problem!

$f(w_j | w_i)$: “Pairs”

- Better solutions?

$(a, *) \rightarrow 32$

Reducer holds this value in memory, rather than the stripe

$(a, b_1) \rightarrow 3$

$(a, b_1) \rightarrow 3 / 32$

$(a, b_2) \rightarrow 12$

$(a, b_2) \rightarrow 12 / 32$

$(a, b_3) \rightarrow 7$

$(a, b_3) \rightarrow 7 / 32$

$(a, b_4) \rightarrow$

$(a, b_4) \rightarrow 1 / 32$

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- The key is to properly sequence data to reducers
 - If it were possible to compute the joint counts before processing the join counts, the reducer could simply divide the joint counts received from mappers by the marginal
 - The notion of “before” and “after” can be captured in the **ordering of key-value pairs**
 - The programmer can define the sort order of keys so that data needed earlier is presented to the reducer before data that is needed later

$f(w_j | w_i)$: “Pairs” – Order Inversion

- A better solution based on order inversion
- The mapper:
 - additionally emits a “special” key of the form $(w_i, *)$
 - The value associated to the special key is one, that represents the contribution
 - Using combiners this will be aggregated before being sent to the reduce
- The reducer:
 - We must make sure that the special key-value pairs are processed before any other key-value pairs where the left word is w_i (define sort order)
 - We also need to guarantee that all pairs associated with the same word are sent to the same reducer (use partitioner)

$f(w_j | w_i)$: “Pairs” – Order Inversion

- Example:
 - The reducer finally receives:

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- The pairs come in order, and thus we can compute the relative frequency immediately.

$f(w_j | w_i)$: “Pairs” – Order Inversion

- Memory requirements:
 - Minimal, because only the marginal (an integer) needs to be stored
 - No buffering of individual co-occurring word
 - No scalability bottleneck
- Key ingredients <https://eduassistpro.github.io/>
 - Emit a special key-value pair to marginal
 - Control the sort order of the intermediate key-value pairs so that the special key-value pair is processed first
 - Define a custom partitioner for routing intermediate key-value pairs

Order Inversion

- Common design pattern
 - Computing relative frequencies requires marginal counts
 - But marginal cannot be computed until you see all counts
 - Buffering is a bad idea!
 - Trick: getting the marginal counts to arrive at the reducer before the joint cou

<https://eduassistpro.github.io/>

- Optimizations
 - Apply in-memory combining pat late marginal counts

Synchronization: Pairs vs. Stripes

- Approach 1: turn synchronization into an ordering problem
 - Sort keys into correct order of computation
 - Partition key space so that each reducer gets the appropriate set of partial results
 - Hold state in reducer across multiple key-value pairs to perform computation
 - Illustrated by the “pairs” approach

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Approach 2: construct data structure to aggregate partial results together
 - Each reducer receives all the data and completes the computation
 - Illustrated by the “stripes” approach

Add WeChat edu_assist_pro

Assignment Project Exam Help

How to Implement Order Inversion

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Implement a Custom Partitioner

- You need to implement a “pair” class first as the key data type
- A customized partitioner extends the *Partitioner* class

```
public static class YourPartitioner extends Partitioner<Key, Value>{
```

- The *key* and *value* are the intermediate key and value produced by the map function

- In the relevant

```
public static class StringPair extends WritableComparable{
```

- It overrides the *getPartition* function

```
public int getPartition(WritableComparable key, Writable value, int numPartitions)
```

- The *numPartitions* is the number of reducers used in the MapReduce program and it is specified in the driver program (by default 1)
- In the relevant frequencies computing problem

```
public int getPartition(StringPair key, Writable value, int numPartitions){  
    return (key.getFirst().hashCode() & Integer.MAX_VALUE) % numPartitions;  
}
```

References

- Chapters 3.3, 3.4, 4.2, 4.3, and 4.4. Data-Intensive Text Processing with MapReduce. Jimmy Lin and Chris Dyer. University of Maryland, College Park.
- Chapter 5 Hadoop I/O. Hadoop The Definitive Guide.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro