

COMP9313: Big Data Management

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 6: Spark

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Assignment Project Exam Help

Part <https://eduassistpro.github.io/> **ction**

Add WeChat edu_assist_pro

Motivation of Spark

- ❑ MapReduce greatly simplified big data analysis on large, unreliable clusters. It is great at one-pass computation.
- ❑ But as soon as it got popular, users wanted more:
 - ❑ More **complex**, multi-pass analytics (e.g. ML, graph)
 - ❑ More **interactive** ad-hoc queries
 - ❑ More **real-time**
- ❑ All 3 need faster processing
 - ❑ One reaction: specialized mode these apps, e.g.,
 - ▶ Pregel (graph processing)
 - ▶ Storm (stream processing)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Limitations of MapReduce

Benefits of data flow: runtime can decide where to run tasks and can automatically recover from failures

Assignment Project Exam Help

- As a general pro<https://eduassistpro.github.io/>
 - It is more suitable for one-pass n a large dataset
 - Hard to compose and nest multi
 - No means of expressing iterative operations
- As implemented in Hadoop
 - All datasets are read from disk, then stored back on to disk
 - All data is (usually) triple-replicated for reliability
 - Not easy to write MapReduce programs using Java

Data Sharing in MapReduce

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Slow due to replication, serialization, and disk IO

- Complex apps, streaming, and interactive queries all need one thing that MapReduce lacks:

Efficient primitives for **data sharing**

Data Sharing in MapReduce

- Iterative jobs involve a lot of disk I/O for each repetition

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- Interactive queries and online processing involves lots of disk I/O



Example: PageRank

- Repeatedly multiply sparse matrix and vector
- Requires repeatedly hashing together page adjacency lists and rank vector

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Hardware for Big Data

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Lots of hard drives of CPUs
Add WeChat edu_assist_pro



And lots of memory!

Goals of Spark

- Keep more data in-memory to improve the performance!
- Extend the MapReduce model to better support two common classes of analytics apps:
 - Iterative algorithms (machine learning, graphs)
 - Interactive data mining
- Enhance program
- Integrate into <https://eduassistpro.github.io/>
- Allow interactive use from Scala

Add WeChat edu_assist_pro

Data Sharing in Spark Using RDD

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

10-100× faster than network and disk

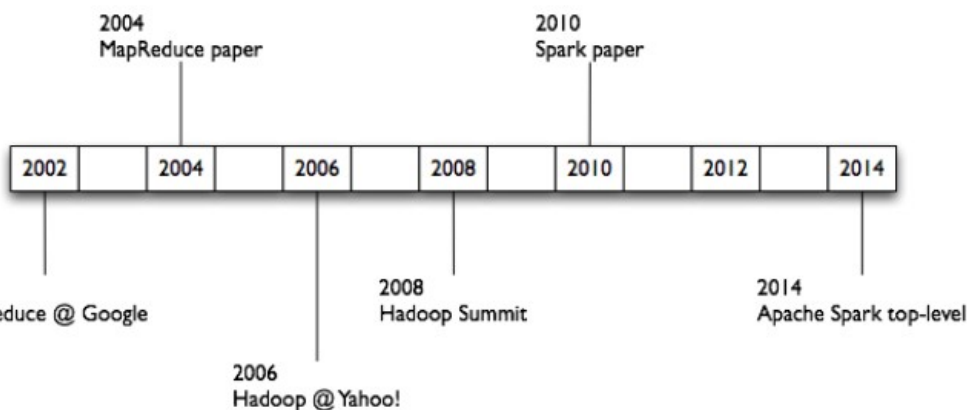
What is Spark

- One popular answer to “What’s beyond MapReduce?”
- Open-source engine for large-scale data processing

- Supports generalized dataflows
- Written in Scala, with bindings in Java and Python

- Brief history

- Developed
- Open-sourc
- Became top-level Apache proje
- Commercial support provided b



What is Spark

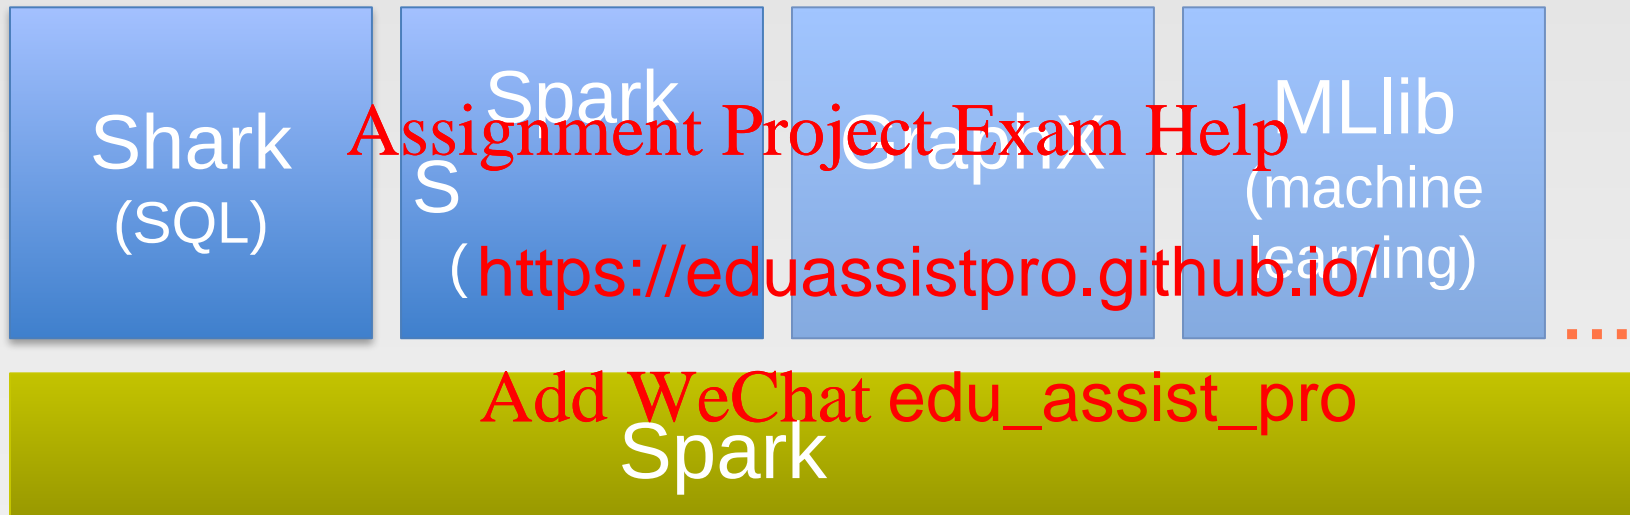
- Fast and expressive cluster computing system interoperable with Apache Hadoop
- Improves efficiency through:
 - **In-memory** computing primitives → Up to 100× faster (10× on disk)
 - General computation graphs
- Improves usability
 - Rich APIs in <https://eduassistpro.github.io/>
 - Interactive shell
- **Spark is not**
 - a modified version of Hadoop
 - dependent on Hadoop because it has its own cluster management
 - Spark uses Hadoop for storage purpose only

Assignment Project Exam Help

Add WeChat → edu_assist_pro 5× less code

What is Spark

- Spark is the basis of a wide set of projects in the Berkeley Data Analytics Stack (BDAS)



- Spark SQL (SQL on Spark)
- Spark Streaming (stream processing)
- GraphX (graph processing)
- MLlib (machine learning library)

Data Sources

- Local Files
 - `file:///opt/httpd/logs/access_log`
- S3
- Hadoop Distributed Filesystem
 - Regular files, sequence files, any other Hadoop InputFormat
- HBase, Cassan

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Spark Ideas

- Expressive computing system, not limited to map-reduce model
- Facilitate system memory
 - avoid saving intermediate results to disk
 - cache data for repetitive queries (e.g. for machine learning)
- Layer an in-memory system on top of Hadoop.
- Achieve fault-tol of replication

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Spark Workflow

- A Spark program first creates a SparkContext object

- Tells Spark how and where to access a cluster

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- Connect to several types of cluster managers (e.g., Mesos, YARN, or its own standalone manager)

- Allocate resources across applications

- Spark executor:

- Run computations
 - Access data storage

Worker Nodes and Executors

- Worker nodes are machines that run executors
 - Host one or multiple Workers
 - One JVM (1 process) per Worker
 - Each Worker can spawn one or more Executors
- Executors run tasks
 - Run in child
 - Execute one <https://eduassistpro.github.io/> ThreadPool

Add WeChat edu_assist_pro

Assignment Project Exam Help

Part <https://eduassistpro.github.io/> **ction**

Add WeChat edu_assist_pro

Scala (Scalable language)

- Scala is a *general-purpose programming language* designed to express common programming patterns in a concise, elegant, and type-safe way
- Scala supports both Object Oriented Programming and Functional Programming
- Scala is Practical
 - Can be use
 - ▶ Mixed S
 - Use existing Java libraries
 - Use existing Java tools (Ant, Maven, ...)
 - Decent IDE Support (NetBeans, IntelliJ, Eclipse)



Why Scala

- Scala supports object-oriented programming. Conceptually, every value is an object and every operation is a method-call. The language supports advanced component architectures through classes and traits
- Scala is also a functional language. Supports functions, immutable data structures and preference for immutability over mutation
- Seamlessly inte
- Being used hea

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Scala Basic Syntax

- When considering a Scala program, it can be defined as a collection of objects that communicate via invoking each other's methods.
- **Object** – same as in Java
- **Class** – same as in Java
- **Methods** – same as in Java
- **Fields** – Each object has one or more instance variables, which are called fields. An instance variable is a variable whose values are assigned to these fields.
- **Traits** – Like Java Interface. A trait is a collection of method and field definitions, which can then be reused by other classes.
- **Closure** – A **closure** is a function, whose return value depends on the value of one or more variables declared outside this function.

closure = function + environment

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Scala is Statically Typed

- You don't have to specify a type in most cases
- Type Inference

```
val sum = 1 + 2 + 3
```

```
val nums = List(1, 2, 3)
```

```
val map = Map("abc" -> List(1, 2, 3))
```

Explicit Types <https://eduassistpro.github.io/>

```
val sum: Int = 1 + 2 + 3
```

```
val nums: List[Int] = List(1
```

```
val map: Map[String, List[Int]] = ...
```

Scala is High level

// Java – Check if string has uppercase character

```
boolean hasUpperCase = false;
```

```
for(int i = 0; i < name.length(); i++) {  
    if(Character.isUpperCase(name.charAt(i))) {  
        hasUpperCase = true;  
        break;  
    }  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

// Scala

```
val hasUpperCase = name.exists(_.isUpper)
```


Scala is Concise

// Java

```
public class Person {  
    private String name;  
    private int age;  
    public Person(String name, Int age) {  
        this.name = name;  
        this.age = age;  
    }  
    public String getName() {  
        return name;  
    }  
    public int getAge() {  
        return age;  
    }  
    public void setName(String name) {  
        this.name = name;  
    }  
    public void setAge(int age) {  
        this.age = age;  
    }  
}
```

// Scala

```
class Person(var name: String, private var _age: Int) {  
    def age = _age // Getter for age  
    def age_=(newAge: Int) { // Setter for age  
        println("Changing age to: "+newAge)  
        _age = newAge  
    }  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Variables and Values

- **Var**iables: values stored can be changed

```
var foo = "foo"  
foo = "bar"  // okay
```

- **Val**ues: immutable variable

```
val foo = "foo"  
foo = "bar"  //
```

Add WeChat edu_assist_pro

Scala is Pure Object Oriented

// Every value is an object

1.toString

// Every operation is a method call

1 + 2 + 3 → (1).+(2).+(3)

// Can omit .toString()

"abc" charAt 1

// Classes (an

```
abstract class Language(val n: String) {  
  override def toString = n  
}
```

// Example implementations

```
class Scala extends Language("Scala")
```

// Anonymous class

```
val scala = new Language("Scala") { /* empty */ }
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Scala Traits

// Like interfaces in Java

```
trait JVM {
```

```
  // But allow implementation
```

```
  override def toString = super.toString + " runs on  
  JVM" }
```

```
trait Static
```

```
  override def  
  Static" }
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

ng+" is

Add WeChat edu_assist_pro

// Traits are stackable

```
class Scala extends Language with JVM with  
  Static {
```

```
  val name = "Scala"
```

```
}
```

```
println(new Scala) → "Scala runs on JVM is Static"
```

Scala is Functional

- First Class Functions. Functions are treated like objects:
 - passing functions as arguments to other functions
 - returning functions as the values from other functions
 - assigning functions to variables or storing them in data structures

Assignment Project Exam Help

```
// Lightweight  
(x:Int) => x + https://eduassistpro.github.io/
```

```
// Calling the anonymous fun  
val plusOne = (x:Int) => x + 1  
plusOne(5) → 6
```

Scala is Functional

- Closures: a function whose return value depends on the value of one or more variables declared outside this function.

// plusFoo can reference any **values/variables** in scope

```
var foo = 1  
val plusFoo = (x:Int
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
plusFoo(5) → 6
```

Add WeChat edu_assist_pro

// Changing foo changes the return value of plusFoo

```
foo = 5
```

```
plusFoo(5) → 10
```

Scala is Functional

□ Higher Order Functions

□ A function that does at least one of the following:

- ▶ takes one or more functions as arguments
- ▶ returns a function as its result

Assignment Project Exam Help

```
val plusOne = (x:In
```

```
val nums = List(1,2
```

```
// map takes a function: Int
```

```
nums.map(plusOne) → Lis
```

```
// Inline Anonymous
```

```
nums.map(x => x + 1) → List(2,3,4)
```

```
// Short form
```

```
nums.map(_ + 1) → List(2,3,4)
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

More Examples on Higher Order Functions

```
val nums = List(1,2,3,4)
```

```
// A few more examples for List class
```

```
nums.exists(_ == 2)           → true
```

```
nums.find(_ == 2)             → Some(2)
```

```
nums.indexOf(2)                → 1
```

```
// functions a https://eduassistpro.github.io/ if true value "1"
```

```
def call(f: Int => Int) = f(1)
```

```
call(plusOne)                 → 2
```

```
call(x => x + 1)              → 2
```

```
call(_ + 1)                   → 2
```


More Examples on Higher Order Functions

```
val basefunc = (x:Int) => ((y:Int) => x + y)
// interpreted by:
basefunc(x){
  sumfunc(y){ return x+y;}
  return sumfunc;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
val closure1 = basefunc(1)           e1(5) = ?
```

Add WeChat edu_assist_pro

```
val closure2 = basefunc(4)           closure2(5) = ?
```

9

- basefunc returns a function, and closure1 and closure2 are of function type.
- While closure1 and closure2 refer to the same function basefunc, the associated environments differ, and the results are different

The Usage of “_” in Scala

□ In anonymous functions, the “_” acts as a placeholder for parameters

```
nums.map(x => x + 1)
```

is equivalent to:

```
nums.map(_ + 1)
```

Assignment Project Exam Help

```
List(1,2,3,4,5).foreach
```

is equivalent to: <https://eduassistpro.github.io/>

```
List(1,2,3,4,5).foreach(a => print(a))
```

Add WeChat edu_assist_pro

□ You can use two or more underscores to refer different parameters.

```
val sum = List(1,2,3,4,5).reduceLeft(_+_)
```

is equivalent to:

```
val sum = List(1,2,3,4,5).reduceLeft((a, b) => a + b)
```

□ The reduceLeft method works by applying the function/operation you give it, and applying it to successive elements in the collection

Assignment Project Exam Help

Par <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Challenge

Existing Systems

- Existing in-memory storage systems have interfaces based on fine-grained updates

- Reads and writes to cells in a table
- E.g., databases, key-value stores, distributed memory

- Requires reads for fault tolerance

-> expensive <https://eduassistpro.github.io/>

- 10-100x slower than memory

Add WeChat edu_assist_pro

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?

Solution: Resilient Distributed Datasets

- *Resilient Distributed Datasets (RDDs)*
 - Distributed collections of objects that can be cached in memory across cluster
 - Manipulated through parallel operators
 - Automatically recomputed on failure based on lineage
- RDDs can express a wide range of operations and capture many current program models
 - Data flow models: MapReduce,
 - Specialized models for iterative

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

What is RDD

- Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. Matei Zaharia, et al. NSDI'12
 - RDD is a **distributed** memory abstraction that lets programmers perform **in-memory** computations on large clusters in a **fault-tolerant** manner.
- Resilient**
 - Fault-tolerant or damaged partitions due to node failure
- Distributed**
 - Data residing on multiple nodes
- Dataset**
 - A collection of partitioned elements, e.g. tuples or other objects (that represent records of the data you work with).
- RDD is the primary data abstraction in Apache Spark and the core of Spark. It enables operations on collection of elements in parallel.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

RDD Traits

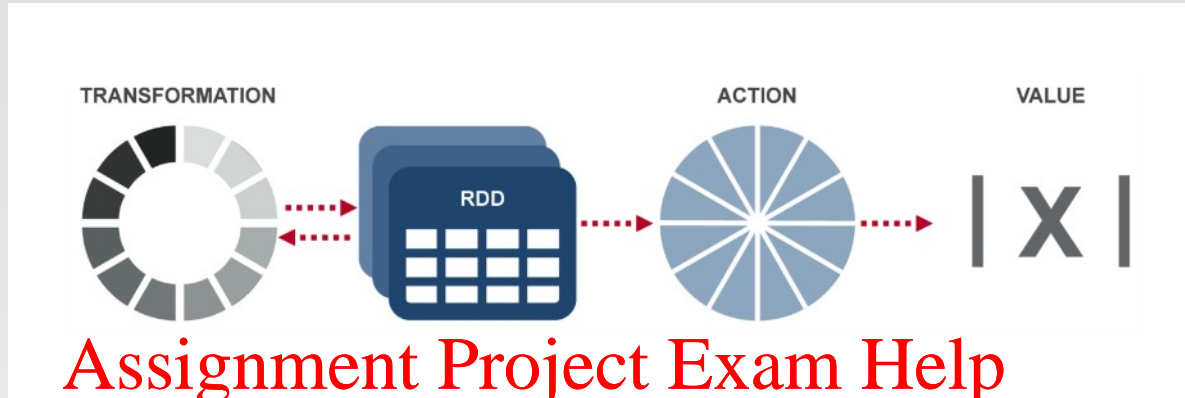
- ❑ **In-Memory**, i.e. data inside RDD is stored in memory as much (size) and long (time) as possible.
- ❑ **Immutable** or **Read-Only**, i.e. it does not change once created and can only be transformed using transformations to new RDDs.
- ❑ **Lazy evaluated**, i.e. the data inside RDD is not available or transformed until an action is executed that triggers the execution.
- ❑ **Cacheable**, i.e. persistent "storage" like memory (default) (the least preferred due to access speed).
- ❑ **Parallel**, i.e. process data in parallel
- ❑ **Typed**, i.e. values in a RDD have types, e.g. RDD[Long] or RDD[(Int, String)].
- ❑ **Partitioned**, i.e. the data inside a RDD is partitioned (split into partitions) and then distributed across nodes in a cluster (one partition per JVM that may or may not correspond to a single node).

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

RDD Operations



Transformation

- Nothing gets written to disk. Transformation function, it just takes an RDD and return a new RDD.
- Transformation functions include *flatMap*, *groupByKey*, *reduceByKey*, *aggregateByKey*, *filter*, *join*, etc.

Action: evaluates and returns a new value.

- When an Action function is called on a RDD object, all the data processing queries are computed at that time and the result value is returned.
- Action operations include *reduce*, *collect*, *count*, *first*, *take*, *countByKey*, *foreach*, *saveAsTextFile*, etc.

Working with RDDs

- Create an RDD from a data source
 - by parallelizing existing collections (lists or arrays)
 - by transforming an existing RDDs
 - from files in HDFS or any other storage system
- Apply transformations to an RDD e.g., map, filter
- Apply actions to

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- Users can control two other aspects:
 - Persistence
 - Partitioning

Creating RDDs

□ From HDFS, text files, Amazon S3, Apache HBase, SequenceFiles, any other Hadoop InputFormat

□ Creating an RDD from a File

□ `val inputfile = sc.textFile("...", 4)`

▶ RDD distributed in 4 partitions

▶ Element

▶ Lazy evaluation opens now

Assignment Project Exam Help

Add WeChat edu_assist_pro

□ Turn a collection into an RDD

□ `sc.parallelize([1, 2, 3])`, creating from a Python list

□ `sc.parallelize(Array("hello", "spark"))`, creating from a Scala Array

□ Creating an RDD from an existing Hadoop InputFormat

□ `sc.hadoopFile(keyClass, valClass, inputFmt, conf)`

Spark Transformations

- Create new datasets from an existing one
- Use lazy evaluation: results not computed right away – instead Spark remembers set of transformations applied to base dataset
 - Spark optimizes the required calculations
 - Spark recovers from failures
- Some transform

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Spark Actions

- Cause Spark to execute recipe to transform source
- Mechanism for getting results out of Spark
- Some action functions

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- Example: `words.collect().foreach(println)`

Example

- Web service is experiencing errors and an operators want to search terabytes of logs in the Hadoop file system to find the cause.

//base RDD

Assignment Project Exam Help

val lines = sc.textFile("hdfs://...")

d RDD

<https://eduassistpro.github.io/>

lines.filter(_.startsWith("Error"))

Add WeChat edu_assist_pro

errors.filter(_.contains("HDFS"))

.map(_.split('\t')(3))

.collect()

- Line1:** RDD backed by an HDFS file (base RDD lines not loaded in memory)
- Line3:** Asks for errors to persist in memory (errors are in RAM)

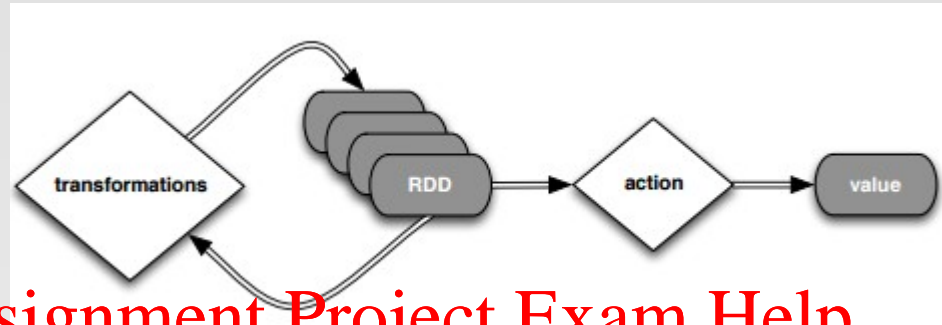
Lineage Graph

- RDDs keep track of *lineage*
- RDD has enough information about how it was derived from to compute its partitions from data in stable storage.

RDD1
Assignment Project Exam Help
RD
<https://eduassistpro.github.io/>
RD
Add WeChat edu_assist_pro
RDD4

- Example:
 - If a partition of errors is lost, Spark rebuilds it by applying a filter on only the corresponding partition of lines.
 - Partitions can be recomputed in parallel on different nodes, without having to roll back the whole program.

Deconstructed



Assignment Project Exam Help

<https://eduassistpro.github.io/>

~~//Transformed RDD~~
Add WeChat **edu_assist_pro**

```
val errors = lines.filter(_.startsWith("Error"))
```

```
errors.persist()
```

```
errors.count()
```

```
errors.filter(_.contains("HDFS"))
```

```
.map(_.split('\t')(3))
```

```
.collect()
```

Deconstructed



//base RDD

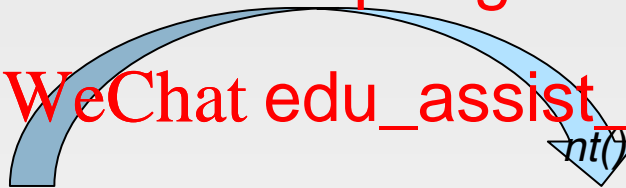
val lines = sc.textFile("hdfs://...")

//Transformed RDD
Assignment Project Exam Help

With("Error"))

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



count() causes Spark to: 1) read data; 2) sum within partitions; 3) combine sums in driver

Put transform and action together:

errors.filter(_.contains("HDFS")).map(_split('\t')(3)).collect()

SparkContext

- SparkContext is the entry point to Spark for a Spark application.
- Once a SparkContext instance is created you can use it to
 - Create RDDs
 - Create accumulators
 - Create broadcast variables
 - access Spark execution environment and acts as the *master application*
- The first thing a Spark program must do is create a SparkContext object, which tells Spark how to access a cluster
- In the Spark shell, a special interpreter-aware SparkContext is already created for you, in the variable called `sc`

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

RDD Persistence: Cache/Persist

- One of the most important capabilities in Spark is *persisting* (or *caching*) a dataset in memory across operations.
- When you persist an RDD, each node stores any partitions of it. You can reuse it in other actions on that dataset
- Each persisted RDD can be stored using a different *storage level*, e.g.
 - MEMORY_ONLY:
 - ▶ Store RDD in the JVM.
 - ▶ If the RDD does not fit in memory, partitions will not be cached and will be recomputed when needed.
 - ▶ This is the default level.
 - MEMORY_AND_DISK:
 - ▶ If the RDD does not fit in memory, store the partitions that don't fit on disk, and read them from there when they're needed.
- `cache() = persist(StorageLevel.MEMORY_ONLY)`

Why Persisting RDD?

```
val lines = sc.textFile("hdfs://...")
```

```
val errors = lines.filter(_.startsWith("Error"))
```

```
errors.persist()
```

```
errors.count()
```

- ❑ If you do `errors.count()` again, the file will be loaded again and computed again
- ❑ Persist will tell Spark to cache the data in memory, to reduce the data loading cost for further actions on the RDD
- ❑ `errors.persist()` will do nothing. It is a no-op. But now the RDD says "read this file and then cache the contents". The action will trigger computation and data caching.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](#)

Spark Key-Value RDDs

- Similar to Map Reduce, Spark supports Key-Value pairs
- Each element of a *Pair RDD* is a pair tuple
- Some Key-Value transformation functions:

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

More Examples on Pair RDD

- Create a pair RDD from existing RDDs

```
val pairs = sc.parallelize( List( ("This", 2), ("is", 3), ("Spark", 5), ("is", 3) ) )  
pairs.collect().foreach(println)
```

Output?

- reduceByKey() function: reduce key-value pairs by key using give *func*

```
val pair1 = pairs.reduceByKey(  
pair1.collect().for
```

<https://eduassistpro.github.io/>

Output?

- mapValues() function: work on value

```
val pair2 = pairs.mapValues( x => x - 1 )  
pairs2.collect().foreach(println)
```

Output?

- groupByKey() function: When called on a dataset of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs

```
pairs.groupByKey().collect().foreach(println)
```

Setting the Level of Parallelism

- All the pair RDD operations take an optional second parameter for number of tasks

```
> words.reduceByKey((x, y) => x + y, 5)
```

```
> words.groupByKey(5)
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Assignment Project Exam Help

Part 4: Setting Model

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

How Spark Works

- User application create RDDs, transform them, and run actions.
 - This results in a DAG (Directed Acyclic Graph) of operators.
 - DAG is compiled into stages
 - Each stage is executed as a series of Task (one Task for each Partition).
- Assignment Project Exam Help**

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

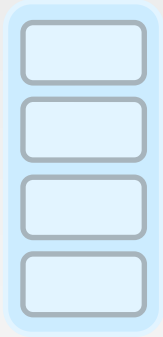
Word Count in Spark

```
val file = sc.textFile("hdfs://...", 4) RDD[String]
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



textFile

Word Count in Spark

```
val file = sc.textFile("hdfs://...", 4)
```

RDD[String]

```
val words = file.flatMap(line =>
```

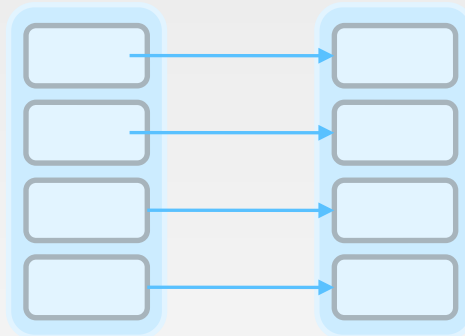
RDD[List[String]]

```
line.split(" "))
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



textFile

flatMap

Word Count in Spark

```
val file = sc.textFile("hdfs://...", 4)
```

RDD[String]

```
val words = file.flatMap(line =>
```

RDD[List[String]]

```
line.split(" "))
```

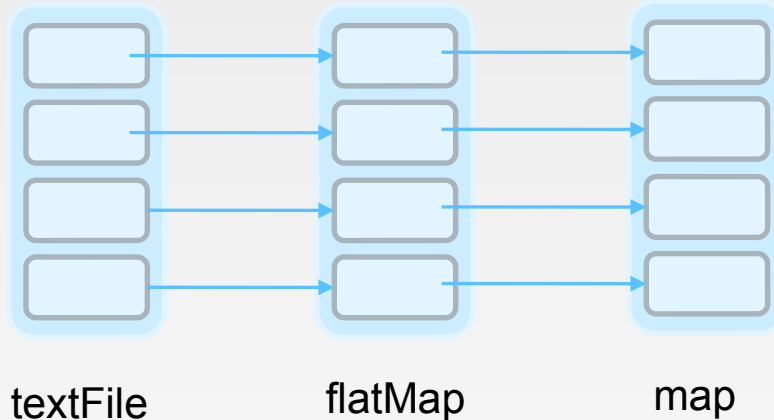
```
val pairs = word
```

RDD[(String, Int)]

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Word Count in Spark

```
val file = sc.textFile("hdfs://...", 4)
```

RDD[String]

```
val words = file.flatMap(line =>
```

RDD[List[String]]

```
line.split(" "))
```

```
val pairs = word
```

RDD[(String, Int)]

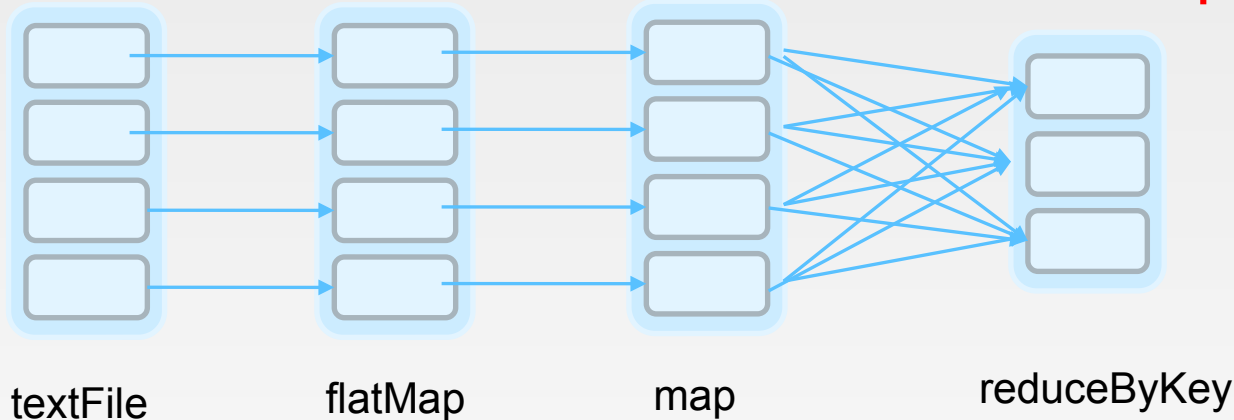
```
val count = pair
```

RDD[(String, Int)]

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Word Count in Spark

```
val file = sc.textFile("hdfs://...", 4)
val words = file.flatMap(line =>
  line.split(" "))
val pairs = word
val count = pair
count.collect()
```

RDD[String]

RDD[List[String]]

RDD[(String, Int)]

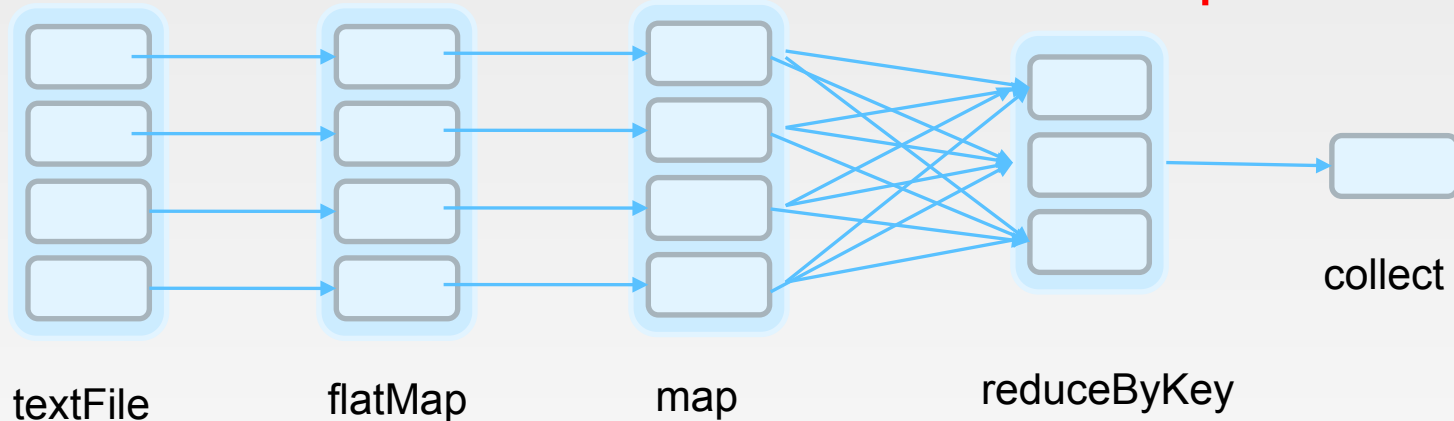
RDD[(String, Int)]

Array[(String, Int)]

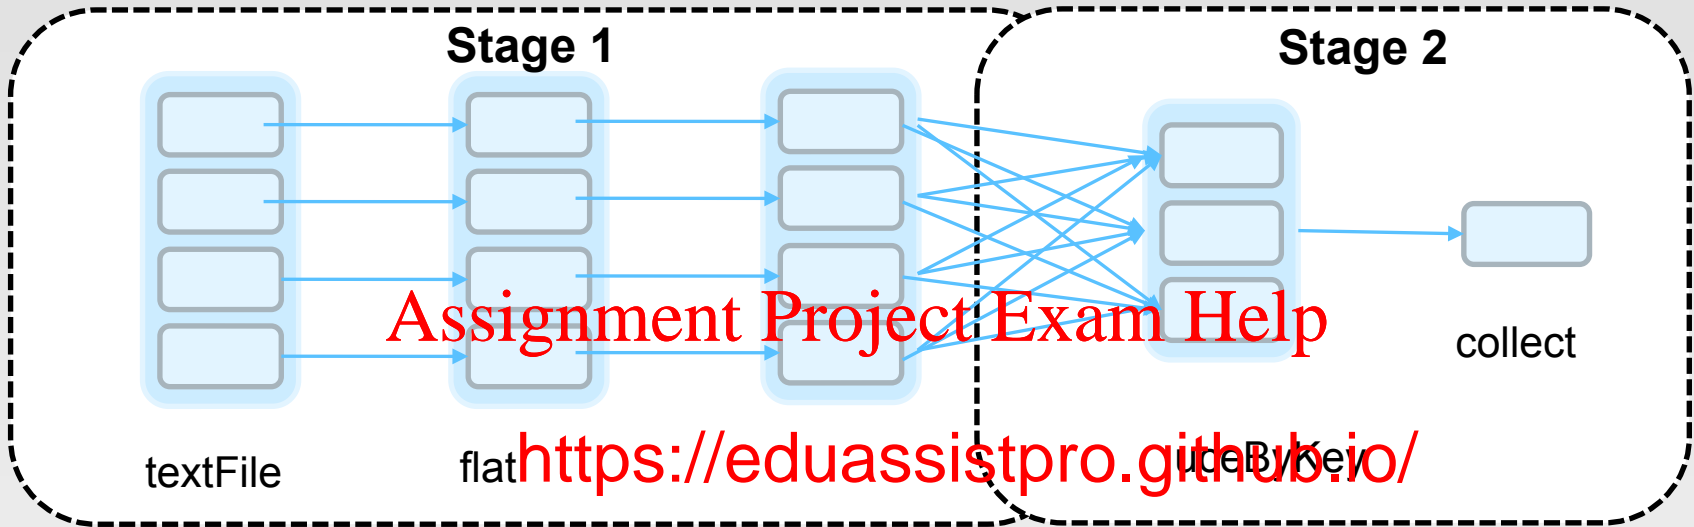
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



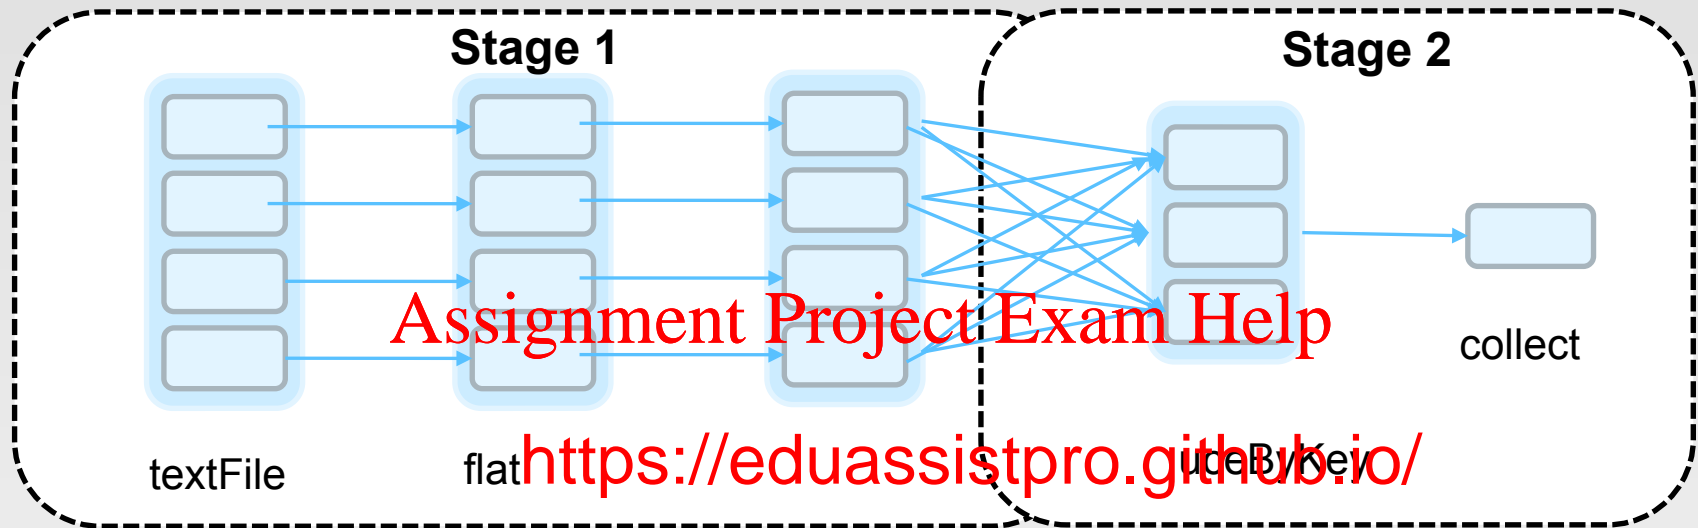
Execution Plan



Add WeChat edu_assist_pro

- The scheduler examines the RDD's lineage graph to build a DAG of stages.
- Stages are sequences of RDDs, that don't have a Shuffle in between
- The boundaries are the shuffle stages.

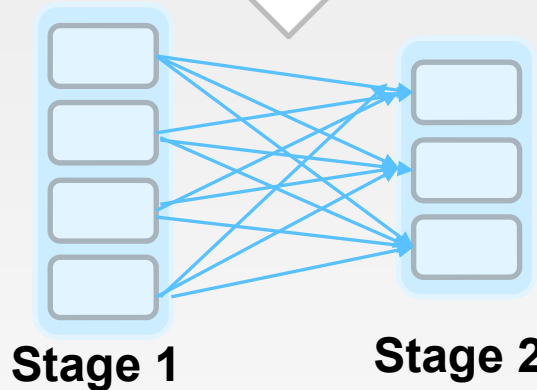
Execution Plan



Add WeChat edu_assist_pro

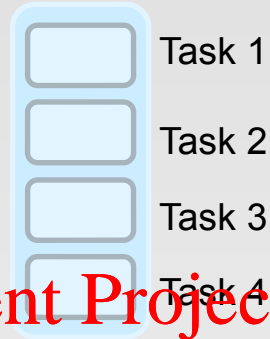


1. Read HDFS split
2. Apply both the maps
3. Start Partial reduce
4. Write shuffle data



1. Read shuffle data
2. Final reduce
3. Send result to driver program

Stage Execution



Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Create a task for each Partition in t
- Serialize the Task
- Schedule and ship Tasks to Slaves
- All this happens internally

Word Count in Spark (As a Whole View)

□ Word Count using Scala in Spark

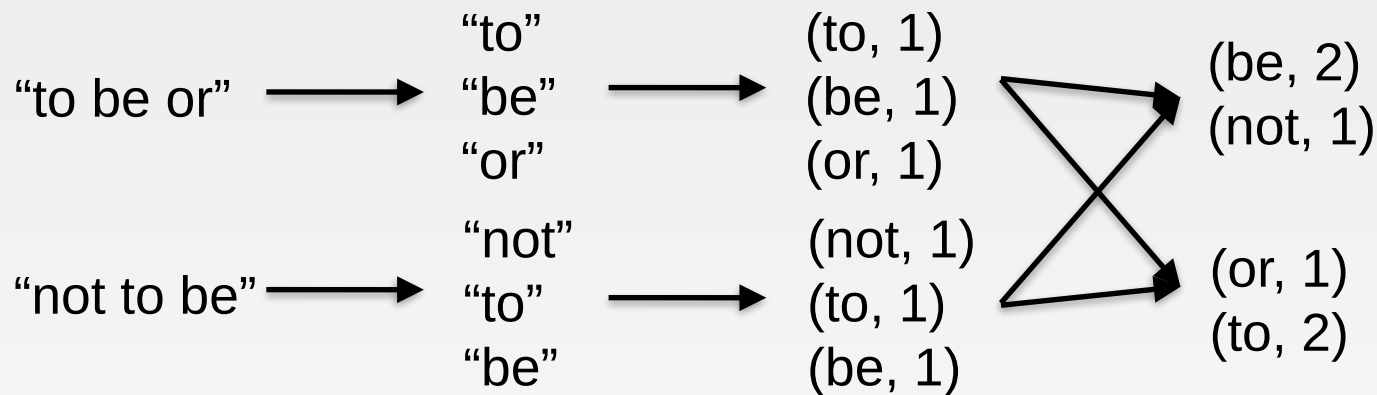
Transformation

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Action

Add WeChat edu_assist_pro



map vs. flatMap

- Sample input file:

```
comp9313@comp9313-VirtualBox:~$ hdfs dfs -cat inputfile
This is a short sentence.
This is a second sentence.
```

```
scala> val inputfile = sc.textFile("inputfile")
inputfile: org.apache.spark.rdd.RDD[String] = inputfile MapPartitionsRDD[1] at t
extFile at <console>:24
```

- map: Return a new RDD by passing each element of the source RDD to the function <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- flatMap: Similar to map, but each input item can be mapped to 0 or more output items (so *func* should return a **Seq** rather than a single item).

```
scala> inputfile.flatMap(x => x.split(" ")).collect()
res4: Array[String] = Array(This, is, a, short, sentence., This, is, a, second, sentence.)
```

RDD Operations

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Spark RDD API Examples:

<http://homepage.cs.latrobe.edu.au/zhe/ZhenHeSparkRDDAPIExamples.html>

Using Local Variables

- Any external variables you use in a closure will automatically be shipped to the cluster:

```
> query = sys.stdin.readline()  
> pages.filter(x => x.contains(query)).count()
```

Assignment Project Exam Help

- Some caveats:

- Each task g <https://eduassistpro.github.io/> sent back)
- Variable must be Serializable

Add WeChat edu_assist_pro

Shared Variables

- When you perform transformations and actions that use functions (e.g., `map(f: T=>U)`), Spark will automatically push a closure containing that function to the workers so that it can run at the workers.

Assignment Project Exam Help

- Any variable or data within a closure or data structure will be distributed to the workers as part of the closure

<https://eduassistpro.github.io/>

- When a function (such as `map` or `reduce`) is executed on a cluster node, it works on separate copies of the data on each node used in it.

Add WeChat edu_assist_pro

- Usually these variables are just constants but they cannot be shared across workers efficiently.

Shared Variables

□ Consider These Use Cases

□ Iterative or single jobs with large global variables

- ▶ Sending large read-only lookup table to workers
- ▶ Sending large feature vector in a ML algorithm to workers
- ▶ Problems? Iterative jobs: need to send large data to each worker with each iter
- ▶ Solution: <https://eduassistpro.github.io/>

□ Counting events that occur during

- ▶ How many input lines were
- ▶ How many input records were corrupt?
- ▶ Problems? Closures are one way: driver -> worker
- ▶ Solution: Accumulators

Broadcast Variables

- Broadcast variables allow the programmer to keep a read-only variable cached on each machine rather than shipping a copy of it with tasks.

- For example, to give every node a copy of a large input dataset efficiently

- Spark also attempts to distribute broadcast variables using efficient broadcast algorithms

- Broadcast variables are created by calling `SparkContext.broadcast(v)`. Its value is accessed by calling the `value` method.

```
scala > val broadcastVar = sc.broadcast(Array(1, 2, 3))
broadcastVar: org.apache.spark.broadcast.Broadcast[Array[Int]] = Broadcast(0)
scala > broadcastVar.value
res0: Array[Int] = Array(1, 2, 3)
```

- The broadcast variable should be used instead of the value `v` in any functions run on the cluster, so that `v` is not shipped to the nodes more than once.

Accumulators

- Accumulators are variables that are only “added” to through an associative and commutative operation and can therefore be efficiently supported in parallel.
- They can be used to implement counters (as in MapReduce) or sums.
- Spark natively supports accumulators of numeric types, and programmers can add support for new types.
- Only driver can r ^{t tasks}
- An accumulator ^{by calling} **SparkContext.accumulator(v).**

```
scala> val accum = sc.longAccumulator("My Accumulator")
accum: org.apache.spark.util.LongAccumulator = LongAccumulator(id: 0, name:
Some(My Accumulator), value: 0)
scala> sc.parallelize(Array(1, 2, 3, 4)).foreach(x => accum.add(x))
... 10/09/29 18:41:08 INFO SparkContext: Tasks finished in 0.317106 s
scala> accum.value
res2: Long = 10
```


Accumulators Example (Python)

- Counting empty lines

```
file = sc.textFile(inputFile)
# Create Accumulator[Int] initialized to 0
blankLines = sc.accumulator(0)
```

```
def extractCallSigns(line):
    global bl
    if (line == " ")
        return line.split(" ")
    else
        return line

callSigns = file.flatMap(extractCallSigns)
print "Blank lines: %d" % blankLines.value
```

- blankLines is created in the driver, and shared among workers
- Each worker can access this variable

References

- <http://spark.apache.org/docs/latest/index.html>
- <http://www.scala-lang.org/documentation/>
- <http://www.scala-lang.org/docu/files/ScalaByExample.pdf>
- [A Brief Intro to Scala](#), by Tim Underwood.
- [Learning Scala, Chapters 1-7.](#)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro